

A FIELD MANUAL FOR THE AGE OF LEVERAGE

ONE PERSON, TEN PEOPLE

*How to become the person who can learn,
build, research and execute almost anything*

RESEARCH → LEARN → THINK → WRITE → BUILD → ANALYSE →
AUTOMATE → SELL

Prabhav Dev

ONE PERSON, TEN PEOPLE

How to become the person who can learn, build, research and execute almost anything

Prabhav Dev

One Person, Ten People

How to become the person who can learn, build, research and execute almost anything

Copyright © 2026 Prabhav Dev. All rights reserved.

No part of this publication may be reproduced, distributed or transmitted in any form or by any means, including photocopying, recording or other electronic or mechanical methods, without the prior written permission of the author, except in the case of brief quotations embodied in critical reviews and certain other non-commercial uses permitted by copyright law.

Disclaimer. This book is sold or distributed for informational and educational purposes only. It is not legal, medical, financial, tax or professional advice, and it is not a substitute for consultation with a qualified professional in the relevant field. The author and publisher make no representations or warranties as to the accuracy, completeness or fitness for a particular purpose of the contents, and specifically disclaim any implied warranties. Techniques, tools and market conditions described here change quickly; capability claims about software systems were accurate to the best of the author's knowledge at the time of writing and may no longer hold. Neither the author nor the publisher shall be liable for any loss of profit or any other commercial or personal damages, including but not limited to special, incidental, consequential or other damages, arising from the use of or reliance upon the material in this book. You are responsible for your own decisions, your own verification, and your own compliance with the laws and professional obligations that apply to you.

Every empirical claim in this book is attributed to a named source in **Sources and Notes** at the back. Where the evidence is contested, thin or self-reported, the text says so. Trade names and trademarks are the property of their respective owners and are used editorially with no claim to ownership and no implication of endorsement.

Set in Source Serif 4, Inter and Inter Tight. Typeset from Markdown with a CSS Paged Media pipeline. 6 × 9 inches.

CONTENTS

<i>The Ten-Person Question</i>	1
PART I • THE GROUND TRUTH	1
1 The Leverage Is Real	1
2 The Ten People Inside One Person	1
3 Cognitive Debt	1
4 Taste Is the Bottleneck	1
PART II • THE LOOP	1
5 Research	1
6 Learn	1
7 Think	1
8 Write	1
9 Build	1
10 Analyse	1
11 Automate	1
12 Sell	1
PART III • THE OPERATING SYSTEM	1
13 The Daily Cadence	1
14 The Sovereign Stack	1
15 The Ninety-Day Plan	1

INTRODUCTION

THE TEN-PERSON QUESTION

In my little group chat with my tech CEO friends there's this betting pool for the first year that there is a one-person billion-dollar company.

— SAM ALTMAN, JANUARY 2024

In April 2026, *The New York Times* published the profile that every person reading this book has now heard some version of. Matthew Gallagher, forty-one, working from his house in Los Angeles, had started a telehealth company with about twenty thousand dollars and roughly two months of effort. He used AI to write the software, produce the advertising creative, draft the website copy, and handle customer service. In 2025, its first full year, the company did \$401 million in sales to around 250,000 customers, at a 16.2 per cent net margin — some \$65 million in profit. For comparison, the reporter noted, the publicly traded incumbent in the same market ran a net margin of 5.5 per cent.

He had one employee. His younger brother.

That is the story as it travelled. Here is the part that did not travel. Six weeks before the profile ran, the US Food and Drug Administration issued the company a warning letter. The agency found that the site displayed compounded drugs bearing the company's own label in a way that falsely suggested the company had compounded them, and that marketing claims implied a federal approval the products did not have. The products were

deemed misbranded. The letter raised the possibility of seizure or injunction. Separately, and less dramatically but more fundamentally, the regulatory window the whole business depended on was already closing: the drug shortages that had made the company's product legal to sell had been formally declared resolved, and the enforcement discretion that followed from them had lapsed the previous year.

THE KEYSTONE

One person used AI to move at extraordinary speed into a narrow, temporary opportunity — and the regulator arrived before the second year was out. Artificial intelligence collapsed the cost of building. It did not collapse the cost of being wrong.

Both halves of that story are true at the same time, and a book that gives you only one of them is not worth your evening. The capability is real. A person working alone now produces, in a week, output that a decade ago required a funded team and a quarter. And the failure mode is also real, and it is not the one people expect. It is not that the work comes out bad. The work comes out *plausible*. It comes out formatted, confident, fluent and wrong in ways that a solo operator — who by definition has nobody to check them — is structurally unequipped to catch.

This book is about closing that gap.

What the title actually claims

One person, ten people is not a promise that you will replace ten salaries. It is a claim about **coverage**.

A ten-person team is not ten times one person. It is a structure in which ten different competences sit in the same room: someone who can find out what is true, someone who can teach the rest, someone who argues against the plan, someone who writes it down, someone who builds it, someone

who reads the numbers, someone who removes the repeated work, someone who sells it, someone who notices the legal exposure, and someone who decides. Most solo operators are excellent at one or two of those and quietly terrible at the rest. The terrible ones are not visible as weaknesses. They are visible as *things that never get done*.

What has changed is that the marginal cost of a competent first attempt at almost any of those functions has collapsed to roughly zero. That is the leverage, and it is genuine. What has not changed — and this is the entire argument of the book — is the cost of knowing whether the attempt was any good.

So the real claim is narrower and more useful than the headline version:

THE KEYSTONE

You can now cover the ground a ten-person team used to cover. You cannot outsource the judgement that decides whether the ground was worth covering, or whether the covering was done correctly. Those two facts define the job.

Where the leverage actually is, which is not where you think

There is now a substantial experimental literature on what AI does to knowledge work, and it does not say what the marketing says. It says something considerably more interesting.

In a pre-registered experiment published in *Science* in 2023, Shakked Noy and Whitney Zhang gave 453 college-educated professionals occupation-specific writing tasks and randomly gave half of them access to ChatGPT. Average time fell by about 40 per cent and blind evaluator-rated quality rose 18 per cent. Very large numbers. Also: the tasks took twenty to thirty minutes, ran to about four hundred words, and — in the authors'

own framing — did not require precise factual accuracy or context about an organisation's goals.

In a field experiment with 758 Boston Consulting Group consultants, a team led by Fabrizio Dell'Acqua and Ethan Mollick found that on eighteen tasks *inside* the model's competence, consultants using GPT-4 completed 12.2 per cent more tasks, 25.1 per cent faster, at more than 40 per cent higher quality. On one task deliberately selected to sit just *outside* that competence, the direction reversed: the control group answered correctly 84.5 per cent of the time and the AI-assisted groups scored 60 and 70.6 per cent. The authors named the underlying shape the **jagged technological frontier** — the model is superhuman at some tasks and confidently incompetent at adjacent ones, and the boundary between them is invisible from the inside.

In a staggered rollout across 5,172 customer-support agents, Erik Brynjolfsson, Danielle Li and Lindsey Raymond measured a 15 per cent average rise in issues resolved per hour. The average conceals the finding. Novice and low-skilled workers gained roughly 34 per cent. The most experienced agents gained almost nothing, and their quality slipped slightly. An agent two months into the job, using the tool, performed like an agent with six months' tenure without it.

And in the result that ought to be printed on the wall of every AI-productivity conference, the research organisation METR took sixteen experienced open-source developers, working on repositories they had contributed to for an average of five years, and randomised 246 real tasks to AI-allowed or AI-disallowed. Allowing AI made them **19 per cent slower**. Before the study, they had forecast a 24 per cent speedup. After living through the slowdown, they still believed they had been sped up by 20 per cent.

Put these together and a clear, non-obvious pattern falls out.

THE KEYSTONE

Your gain from AI is largest exactly where you are weakest, and smallest — possibly negative — in the one domain where you are genuinely expert. The novice gains a third. The expert gains nothing and cannot tell.

That is not a disappointing finding. For the purposes of this book it is the *whole* finding, because it is precisely the empirical case for “one person, ten people.” You do not need AI to help you with the thing you are already good at. You need it for the nine roles you never trained for. That is where the measured effect sizes are enormous. The argument for breadth is stronger than the argument for speed, and it is better supported by evidence.

One more result, for honesty’s sake. Anders Humlum and Emilie Vestergaard linked adoption surveys covering some 25,000 workers across eleven exposed occupations in Denmark to administrative employer–employee records. Two years after ChatGPT’s launch, the effect on earnings and hours was a precise zero — the estimates rule out effects larger than 2 per cent. Workers reported saving about 3 per cent of their time, and almost none of it reached their pay.

Nothing about having a powerful tool automatically converts into value. The conversion is the work. It is what this book is for.

The loop

The rest of this book is organised around eight activities. Together they describe essentially everything a knowledge business does, and they run in a cycle: what you learn at the end of one turn changes what you research at the start of the next.



FIGURE 1 The loop. Eight stations, one operator. Each station is a place where AI can do a large share of the production and none of the judging. The direction of travel matters less than the fact that it closes — the analysis of what you sold is the research for what you build next.

Read the figure as a working diagram, not a poster. Four things about it are load-bearing.

It is a cycle, not a funnel. Most frameworks of this shape are pipelines that end in a sale. This one ends where it started, because the information that comes back from selling is the highest-quality research input you will ever get, and it is the one input your competitors cannot copy.

You can enter at any station. Nobody actually starts at Research. You start wherever the live problem is. The value of the map is not the order; it is the reminder that the other seven stations exist and that you are probably neglecting three of them.

Each station has the same internal structure. There is a part that is production — generating, drafting, scaffolding, summarising, transforming — and a part that is judgement — deciding what is worth doing, and deciding whether what came back is true and good. AI has become extremely strong at the first and remains unreliable at the second. Every chapter in Part II works the same seam: here is what to hand over, here is what you must keep, and here is the specific mechanism for checking the handover.

You are in the middle for a reason. In a ten-person team, quality control is a by-product of structure: someone else reads the brief, someone else reviews the pull request, someone else asks why in the meeting. Alone, you have no structural redundancy. You have to *manufacture* the second opinion, deliberately, as a habit. That is the single most important operational difference between working alone with AI and working on a team, and it is the thread running through every chapter.

What this book will not do

It will not tell you that you can skip learning things. The evidence points hard the other way, and Chapter 3 is devoted to it. In a controlled study at MIT's Media Lab, participants who wrote essays with an AI assistant showed measurably weaker recall of what they had just produced and

weaker neural signatures of engagement, and the effects persisted when the assistance was withdrawn. In a larger educational study, learners using ChatGPT produced *better essays* while showing *no greater knowledge gain* than learners without it. Those two findings together describe the central hazard of this entire way of working: the artefact improves while the person does not, and the artefact is the only thing you can see.

It will not give you a tool list that expires. Specific products change every few months. The distinction between a task whose output you can verify cheaply and one whose output you cannot has not changed in fifty years of automation research and will outlive every model named in this book. Where a tool is named, it is named as an example, and the reasoning is written so you can re-derive the choice when the tool is gone.

It will not pretend the research is settled. Some of what follows rests on large pre-registered experiments. Some rests on a single well-documented case. Some rests on practitioner consensus that has never been tested properly. The text tells you which is which, every time, because a book about verification that asks you to take its own claims on trust would be a bad joke. Every empirical claim is sourced at the back.

And it will not promise you a billion-dollar company. The “one-person unicorn” line began, as the epigraph records, as a bet in a group chat. As of this writing no single-founder company has reached a verified billion-dollar valuation. The defensible version of the claim is about revenue per person, not valuation; about two or three people, not one; and about leverage as a direction of travel, not a destination. That version is still remarkable. It does not need inflating.

How to read this

Part I: The Ground Truth establishes what actually changed, what did not, and what it costs you. Chapter 1 is the evidence on where the leverage is. Chapter 2 breaks the ten-person team into its actual functions and has you audit yourself against them. Chapter 3 is the bill — the cognitive debt

you accrue by offloading, and how to avoid paying it. Chapter 4 is about the thing that becomes scarce when production becomes free.

Part II: The Loop is eight chapters, one per station. Each follows the same architecture — what the station is for, what the evidence says, what to delegate, what to keep, the failure mode, a drill, and the prompts. These are reference chapters. Read them in order once, then return to them out of order forever.

Part III: Running the System turns the map into a calendar. Chapter 13 is the weekly cadence. Chapter 14 is the catalogue of ways this goes wrong, including the regulatory one that opened this introduction. Chapter 15 is a ninety-day plan to go from reading about it to operating it.

Then the appendices, which are the parts you will actually use most: the prompt library organised by station, the verification protocol, and the ten-person audit worksheet.

BEFORE YOU GO ON

Write down, in one sentence, the last piece of work you shipped that you could not have produced alone three years ago. Then write down how you knew it was correct. If the second sentence is harder than the first, you are the reader this book was written for.

One person can now do the work of ten. Whether that produces ten times the value or ten times the volume of confident, well-formatted error depends entirely on what you do with the part that cannot be delegated. That part is you. Let us make it good.

PREVIEW

This is the end of the sample.

You have the cover, the contents, and the introduction. The rest of the book stays behind checkout.

The complete 177-page PDF includes:

- Fifteen chapters across three parts
- The eight-station loop, Research through Sell
- The prompt library and verification playbook
- Worksheets and the 90-day plan

[payhip.prabhavdev.com](https://payhip.com/prabhavdev)

Prabhav Dev